

Aim: To install Kubectl and execute Kubectl commands to manage the Kubernetes cluster and deploy Your First Kubernetes Application.

1. Create an EC2 instance with OS as Amazon Linux and make sure to allow SSH traffic.

| Name | Instance ID | Instance state | Instance type | Status check | Alarm status | Availability Zone | Public IPv4 DNS | Public IP |
|--------|---------------------|----------------|---------------|--------------|--------------|-------------------|-------------------------|------------|
| kubern | i-0c873b5b4ac7c0767 | Running | t2.medium | Initializing | View alarms | us-east-1a | ec2-3-91-46-250.comp... | 3.91.10.10 |

2. To install docker run the following command:
sudo yum install docker -y

```

created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.
Verifying      : containerd-1.7.20-1.amzn2023.0.1.x86_64                                1/10
Verifying      : docker-25.0.6-1.amzn2023.0.2.x86_64                                2/10
Verifying      : iptables-libs-1.8.8-3.amzn2023.0.2.x86_64                           3/10
Verifying      : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64                           4/10
Verifying      : libcgrou-3.0-1.amzn2023.0.1.x86_64                                 5/10
Verifying      : libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64                  6/10
Verifying      : libnftnl-1.0.1-19.amzn2023.0.2.x86_64                             7/10
Verifying      : libnftnl-1.2.2-2.amzn2023.0.2.x86_64                             8/10
Verifying      : pigz-2.5-1.amzn2023.0.3.x86_64                                    9/10
Verifying      : runc-1.1.13-1.amzn2023.0.1.x86_64                                10/10
Installed:
containerd-1.7.20-1.amzn2023.0.1.x86_64  docker-25.0.6-1.amzn2023.0.2.x86_64  iptables-libs-1.8.8-3.amzn2023.0.2.x86_64  iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
libcgrou-3.0-1.amzn2023.0.1.x86_64        libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64  libnftnl-1.0.1-19.amzn2023.0.2.x86_64  libnftnl-1.2.2-2.amzn2023.0.2.x86_64
pigz-2.5-1.amzn2023.0.3.x86_64            runc-1.1.13-1.amzn2023.0.1.x86_64
Complete!

```

3. Configure cgroup in daemon.json file using the following commands:

```

cd /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF

```

```
[ec2-user@ip-172-31-27-89 ~]$ cd /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
```

4. Run the following command after this:

```
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
```

```
[ec2-user@ip-172-31-27-89 docker]$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[ec2-user@ip-172-31-27-89 docker]$
```

5. Verify Installation

```
[ec2-user@ip-172-31-27-89 docker]$ docker -v
Docker version 25.0.5, build 5dc9bcc
```

6. Install Kubernetes

Disable SELinux before configuring kubelet

```
sudo setenforce 0
```

```
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/'
/etc/selinux/config
```

Add kubernetes repository

```
cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
```

```

gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.k
ey
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF

```

```

[ec2-user@ip-172-31-27-89 docker]$ sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
[ec2-user@ip-172-31-27-89 docker]$ cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
[ec2-user@ip-172-31-27-89 docker]$

```

```

sudo yum update
sudo yum install -y kubelet kubeadm kubectl
--disableexcludes=kubernetes

```

```

[ec2-user@ip-172-31-27-89 docker]$ sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Kubernetes
Dependencies resolved.
59 kB/s | 9.4 kB    00:00

```

| Package | Architecture | Version | Repository | Size |
|---------------------------------|--------------|-----------------------|-------------|-------|
| Installing: | | | | |
| kubeadm | x86_64 | 1.31.1-150500.1.1 | kubernetes | 11 M |
| kubectl | x86_64 | 1.31.1-150500.1.1 | kubernetes | 11 M |
| kubelet | x86_64 | 1.31.1-150500.1.1 | kubernetes | 15 M |
| Installing dependencies: | | | | |
| conntrack-tools | x86_64 | 1.4.6-2.amzn2023.0.2 | amazonlinux | 208 k |
| cri-tools | x86_64 | 1.31.1-150500.1.1 | kubernetes | 6.9 M |
| kubernetes-cni | x86_64 | 1.5.1-150500.1.1 | kubernetes | 7.1 M |
| libnetfilter_cthelper | x86_64 | 1.0.0-21.amzn2023.0.2 | amazonlinux | 24 k |
| libnetfilter_cttimeout | x86_64 | 1.0.0-19.amzn2023.0.2 | amazonlinux | 24 k |
| libnetfilter_queue | x86_64 | 1.0.5-2.amzn2023.0.2 | amazonlinux | 30 k |

```

Transaction Summary

```

Configure internet options to allow bridging

- sudo swapoff -a
- echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
- sudo sysctl -p

```

[ec2-user@ip-172-31-27-89 docker]$ sudo swapoff -a
echo "net.bridge.bridge-nf-call-iptables=1" | sudo tee -a /etc/sysctl.conf
sudo sysctl -p

net.bridge.bridge-nf-call-iptables=1
net.bridge.bridge-nf-call-iptables = 1
[ec2-user@ip-172-31-27-89 docker]$

```

7. Initialize the kubecoluster

```
sudo kubeadm init --pod-network-cidr=10.244.0.0/16
```

```
[ec2-user@ip-172-31-27-89 docker]$ kubectl apply -f https://k8s.io/examples/pods/simple-pod.yaml
pod/nginx created
[ec2-user@ip-172-31-27-89 docker]$
```

Run 3 commands

```
[ec2-user@ip-172-31-27-89 docker]$ mkdir -p $HOME/.kube && sudo cp -f /etc/kubernetes/admin.conf $HOME/.kube/config && sudo chown $(id -u):$(id -g) $HOME/.kube/config
cp: cannot stat '/etc/kubernetes/admin.conf': No such file or directory
[ec2-user@ip-172-31-27-89 docker]$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[WARNING FileExisting-socat]: socat not found in system path
[WARNING FileExisting-tc]: tc not found in system path
[WARNING Service-Kubelet]: kubelet service is not enabled, please run 'systemctl enable kubelet.service'
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
```

8. Deploy nginx server on this cluster using the command

kubectl apply -f <https://k8s.io/examples/pods/simple-pod.yaml>

Also run kubectl get pods to check creation of pod

```
[ec2-user@ip-172-31-27-89 docker]$ kubectl apply -f https://k8s.io/examples/pods/simple-pod.yaml
pod/nginx created
[ec2-user@ip-172-31-27-89 docker]$ kubectl describe pod nginx
Name:          nginx
Namespace:     default
Priority:       0
Service Account: default
Node:          <none>
Labels:        <none>
Annotations:   <none>
Status:        Pending
IP:            <none>
IPs:           <none>
Containers:
  nginx:
    Image:      nginx:1.14.2
    Port:       80/TCP
    Host Port:  0/TCP
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-sbq47 (ro)
Conditions:
  Type          Status
```

9. Use kubectl get pods to check status of pods

10. Also mention the port that u want to host

```
Forwarding from 127.0.0.1:8081 -> 80
Forwarding from [::1]:8081 -> 80
```

Conclusion:

In this experiment, we set up a Kubernetes cluster on an Amazon Linux EC2 instance. We installed Docker, configured it to use systemd for cgroup management, and prepared Kubernetes by disabling SELinux and installing necessary components. After initializing the cluster and deploying the Flannel network plugin, we successfully launched an Nginx server. We ensured the Nginx pod was accessible on port from the local machine.